

Multimodal Classification Pipeline for ECG Data using ML and LSTM Models

Import Libraries:

```
[ ]: # Data manipulation and visualization
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Scikit-learn: preprocessing, decomposition, splitting, and metrics
from sklearn.preprocessing import MinMaxScaler, label_binarize
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (accuracy_score, classification_report,
    ↪confusion_matrix,
                                roc_auc_score, roc_curve, auc,
    ↪precision_recall_curve,
                                average_precision_score)

# TensorFlow/Keras for deep learning
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, LSTM, Dropout
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.utils import to_categorical

# IPython for DataFrame display
from IPython.display import display

sns.set(style="whitegrid")
plt.rcParams['figure.figsize'] = (10, 6)
```

Data Loading:

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

```

# Define file paths
train_path = "/content/drive/My Drive/ECG_Dataset/mitbih_train.csv"

# Load CSV files
df_train = pd.read_csv(train_path, header=None)

# Display first few rows of dataset
print("First few rows of the training dataset:")
display(df_train.head())

```

Mounted at /content/drive

First few rows of the training dataset:

	0	1	2	3	4	5	6	\
0	0.977941	0.926471	0.681373	0.245098	0.154412	0.191176	0.151961	
1	0.960114	0.863248	0.461538	0.196581	0.094017	0.125356	0.099715	
2	1.000000	0.659459	0.186486	0.070270	0.070270	0.059459	0.056757	
3	0.925414	0.665746	0.541436	0.276243	0.196133	0.077348	0.071823	
4	0.967136	1.000000	0.830986	0.586854	0.356808	0.248826	0.145540	

	7	8	9	...	178	179	180	181	182	183	184	185	\
0	0.085784	0.058824	0.049020	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
1	0.088319	0.074074	0.082621	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
2	0.043243	0.054054	0.045946	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
3	0.060773	0.066298	0.058011	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
4	0.089202	0.117371	0.150235	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

	186	187
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0
3	0.0	0.0
4	0.0	0.0

[5 rows x 188 columns]

Categorization of Data:

```

[ ]: # --- Missing Values ---
# missing values in data
missing_train = pd.DataFrame(df_train.isnull().sum(), columns=["Missing_
↵Values"])

print("Missing values in training data:")
display(missing_train)

```

Missing values in training data:

	Missing Values
0	0

```

1          0
2          0
3          0
4          0
..         ...
183        0
184        0
185        0
186        0
187        0

```

[188 rows x 1 columns]

```

[ ]: # --- Data Types ---
dtypes_train = pd.DataFrame(df_train.dtypes, columns=["Data Type"])
print("\nTraining data types:")
display(dtypes_train)

```

Training data types:

```

      Data Type
0      float64
1      float64
2      float64
3      float64
4      float64
..         ...
183     float64
184     float64
185     float64
186     float64
187     float64

```

[188 rows x 1 columns]

```

[ ]: # --- Descriptive Statistics ---
desc_train = df_train.describe().T
print("\nDescriptive Statistics of Training Data:")
display(desc_train)

```

Descriptive Statistics of Training Data:

	count	mean	std	min	25%	50%	75%	max
0	87554.0	0.890360	0.240909	0.0	0.921922	0.991342	1.000000	1.0
1	87554.0	0.758160	0.221813	0.0	0.682486	0.826013	0.910506	1.0
2	87554.0	0.423972	0.227305	0.0	0.250969	0.429472	0.578767	1.0
3	87554.0	0.219104	0.206878	0.0	0.048458	0.166000	0.341727	1.0
4	87554.0	0.201127	0.177058	0.0	0.082329	0.147878	0.258993	1.0

183	87554.0	0.003471	0.036255	0.0	0.000000	0.000000	0.000000	1.0
184	87554.0	0.003221	0.034789	0.0	0.000000	0.000000	0.000000	1.0
185	87554.0	0.002945	0.032865	0.0	0.000000	0.000000	0.000000	1.0
186	87554.0	0.002807	0.031924	0.0	0.000000	0.000000	0.000000	1.0
187	87554.0	0.473376	1.143184	0.0	0.000000	0.000000	0.000000	4.0

[188 rows x 8 columns]

Visualize Class Distribution:

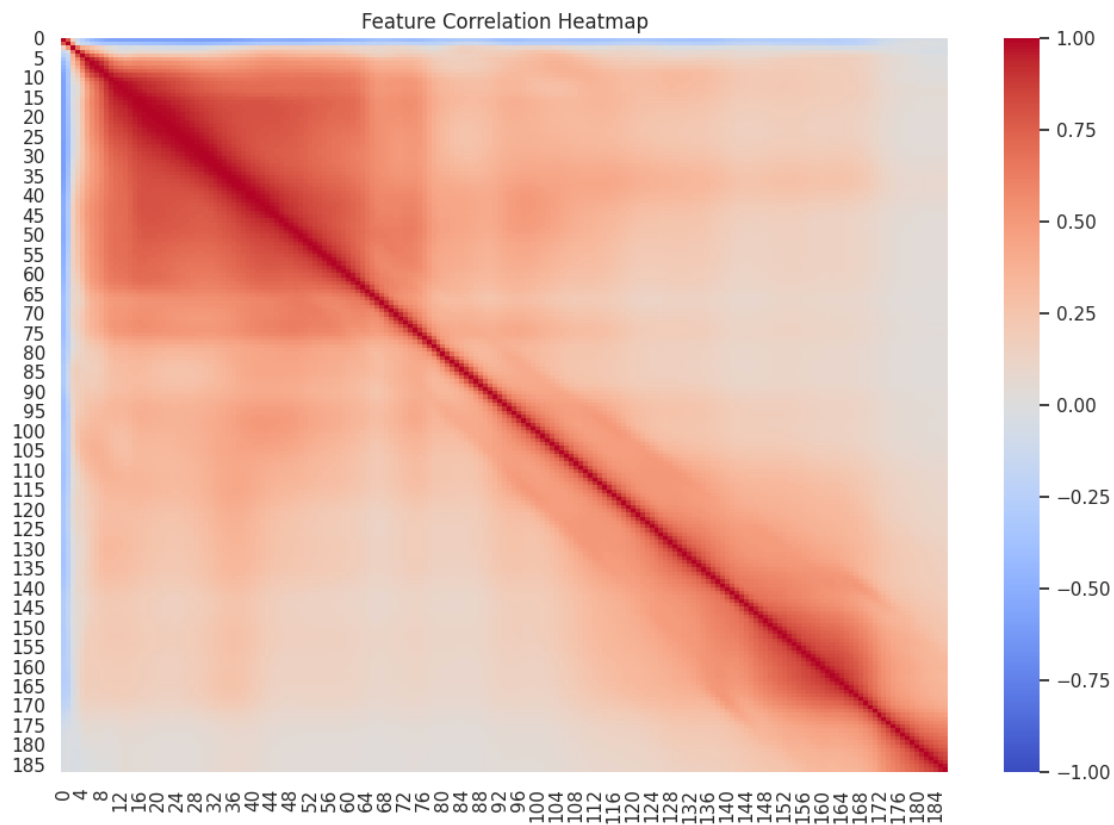
```
[ ]: # Class distribution
plt.figure()
sns.countplot(x=df_train.iloc[:, -1])
plt.title("Class Distribution in Training Data")
plt.xlabel("Class Labels")
plt.ylabel("Count")
plt.show()
```



Compute and Visualize the Correlation Matrix:

```
[ ]: # Correlation matrix
correlation_matrix = df_train.iloc[:, :-1].corr()
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, cmap="coolwarm", vmin=-1, vmax=1, annot=False)
```

```
plt.title("Feature Correlation Heatmap")
plt.show()
```

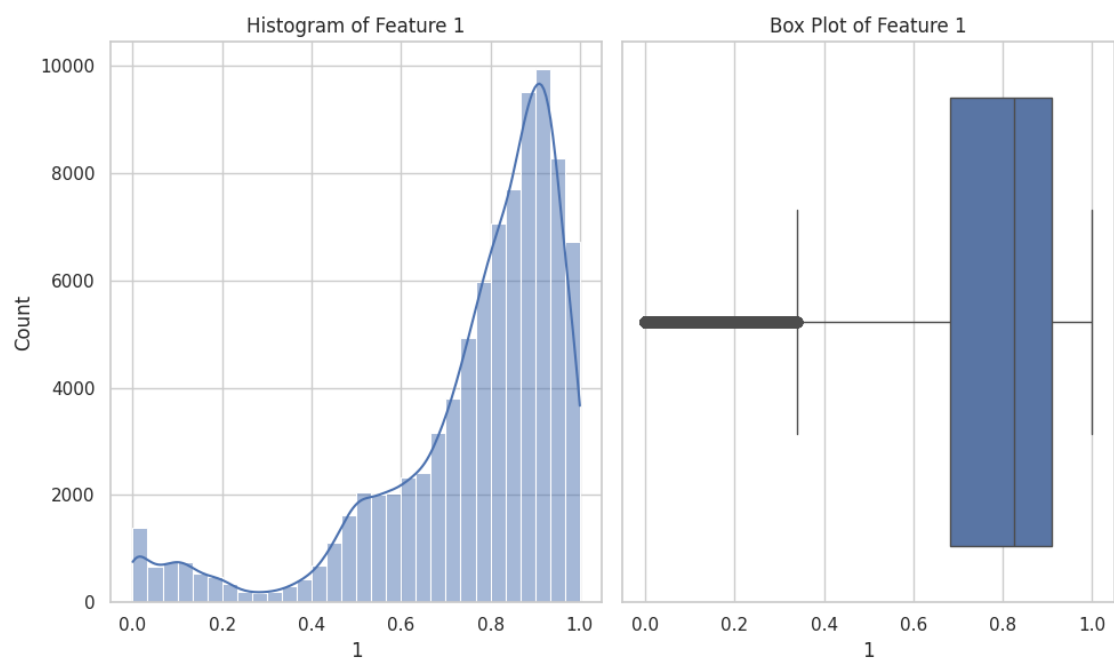




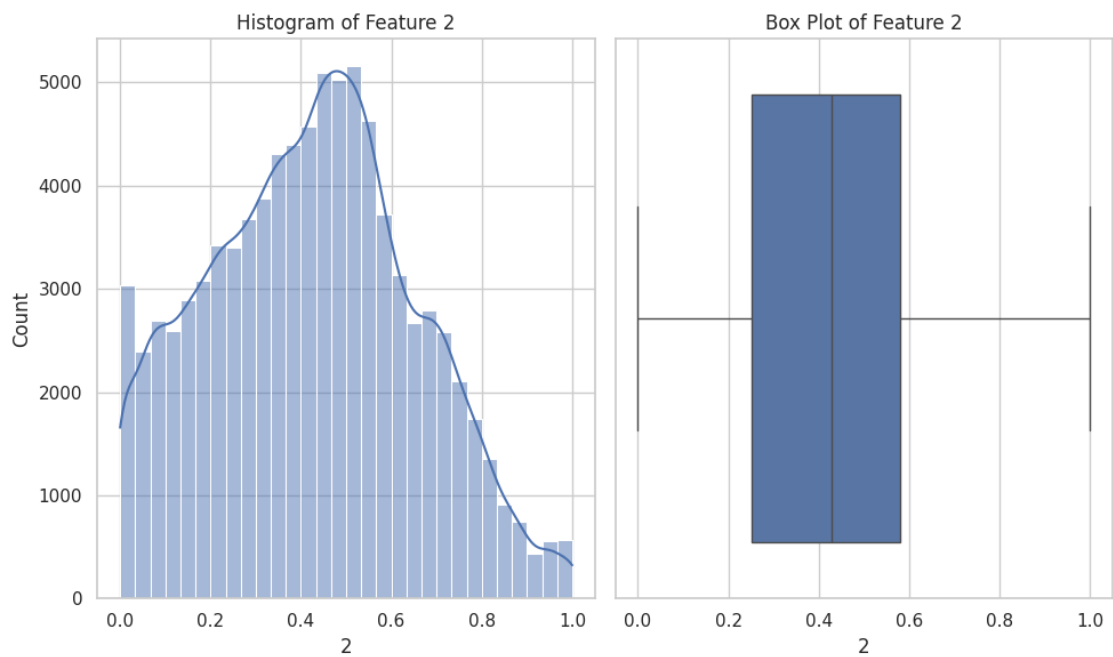
Additional Visualizations for Feature Distributions:


```
[ ]: selected_features = [0, 1, 2, 3, 4]
for feature in selected_features:
    plt.figure()
    plt.subplot(1, 2, 1)
    sns.histplot(df_train[feature], bins=30, kde=True)
    plt.title(f"Histogram of Feature {feature}")

    plt.subplot(1, 2, 2)
    sns.boxplot(x=df_train[feature])
    plt.title(f"Box Plot of Feature {feature}")

plt.tight_layout()
plt.show()
```







Data Cleaning

```
[ ]: # Remove duplicate rows
initial_rows = df_train.shape[0]
df_train = df_train.drop_duplicates()
final_rows = df_train.shape[0]
print(f"Removed {initial_rows - final_rows} duplicate rows from training data.")

def remove_outliers(df, column, factor=1.5):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - factor * IQR
    upper_bound = Q3 + factor * IQR
    return df[(df[column] >= lower_bound) & (df[column] <= upper_bound)]

df_train = remove_outliers(df_train, 0)
print(f"Training data shape after outlier removal: {df_train.shape}")
```

Removed 0 duplicate rows from training data.

Training data shape after outlier removal: (75393, 188)

Data Scaling(Normalization)

```
[ ]: # Separate features and target from raw DataFrames
X_train = df_train.iloc[:, :-1]
```



```

y_train = df_train.iloc[:, -1]

# Normalize features using MinMaxScaler
scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)

df_train_scaled = pd.DataFrame(X_train_scaled, columns=X_train.columns)
df_train_scaled['target'] = y_train.values

print("First few rows of the normalized training data:")
display(df_train_scaled.head())

```

First few rows of the normalized training data:

	0	1	2	3	4	5	6	\
0	0.977941	0.926471	0.681373	0.245098	0.154412	0.191176	0.151961	
1	0.960114	0.863248	0.461538	0.196581	0.094017	0.125356	0.099715	
2	1.000000	0.659459	0.186486	0.070270	0.070270	0.059459	0.056757	
3	0.925414	0.665746	0.541436	0.276243	0.196133	0.077348	0.071823	
4	0.967136	1.000000	0.830986	0.586854	0.356808	0.248826	0.145540	

	7	8	9	...	178	179	180	181	182	183	184	185	\
0	0.085784	0.058824	0.049020	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
1	0.088319	0.074074	0.082621	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
2	0.043243	0.054054	0.045946	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
3	0.060773	0.066298	0.058011	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
4	0.089202	0.117371	0.150235	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

	186	target
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0
3	0.0	0.0
4	0.0	0.0

[5 rows x 188 columns]

Feature Engineering

```

[ ]: # Apply PCA to reduce dimensionality
pca = PCA(n_components=50)
X_train_pca = pca.fit_transform(X_train_scaled)

df_train_pca = pd.DataFrame(X_train_pca, columns=[f"PC{i+1}" for i in_
↪range(X_train_pca.shape[1])])
df_train_pca['target'] = y_train.values

print("Shape of PCA-transformed training data:", df_train_pca.shape)

```

Shape of PCA-transformed training data: (87554, 51)

```
[ ]: # Use normalized data for modeling
X_train_full = df_train_scaled.drop('target', axis=1)
y_train_full = df_train_scaled['target']

# Split into training (80%) and validation (20%) sets
X_train_split, X_val_split, y_train_split, y_val_split = train_test_split(
    X_train_full, y_train_full, test_size=0.2, random_state=42,
    ↪stratify=y_train_full
)

print("Training set shape:", X_train_split.shape)
print("Validation set shape:", X_val_split.shape)
```

Training set shape: (70043, 187)

Validation set shape: (17511, 187)

Logistic Regression with GridSearchCV

```
[ ]: param_grid_lr = {
    'C': [0.01, 0.1, 1, 10],
    'solver': ['liblinear', 'saga'],
    'penalty': ['l2']
}

lr = LogisticRegression(max_iter=1000)
grid_lr = GridSearchCV(lr, param_grid_lr, cv=5, scoring='accuracy', n_jobs=-1)
grid_lr.fit(X_train_split, y_train_split)

print("Best parameters for Logistic Regression:", grid_lr.best_params_)
y_val_pred_lr = grid_lr.predict(X_val_split)
accuracy_lr = accuracy_score(y_val_split, y_val_pred_lr)
print("Validation Accuracy for Logistic Regression:", accuracy_lr)
```

Best parameters for Logistic Regression: {'C': 10, 'penalty': 'l2', 'solver': 'saga'}

Validation Accuracy for Logistic Regression: 0.9530472843026726

Random Forest

```
[ ]: # Initialize Random Forest Classifier
rf_model = RandomForestClassifier(n_estimators=100, class_weight="balanced",
    ↪random_state=42)

rf_model.fit(X_train_split, y_train_split)

# Predict on validation data
```

```

y_pred_rf = rf_model.predict(X_val_split)

print("Random Forest Performance:")
print(classification_report(y_val_split, y_pred_rf))

```

Random Forest Performance:

	precision	recall	f1-score	support
0.0	0.98	1.00	0.99	13122
1.0	0.98	0.63	0.77	390
2.0	0.98	0.84	0.90	596
3.0	0.81	0.62	0.70	112
4.0	1.00	0.94	0.97	859
accuracy			0.98	15079
macro avg	0.95	0.80	0.86	15079
weighted avg	0.98	0.98	0.97	15079

Data prepration for LSTM

```

[ ]: n_timesteps = X_train_full.shape[1]

# Reshape for LSTM: (samples, timesteps, features)
X_train_lstm = X_train_split.values.reshape(-1, n_timesteps, 1)
X_val_lstm = X_val_split.values.reshape(-1, n_timesteps, 1)

num_classes = len(np.unique(y_train_full))
y_train_cat = to_categorical(y_train_split, num_classes)
y_val_cat = to_categorical(y_val_split, num_classes)

```

LSTM with Early Stopping

```

[ ]: # Building LSTM model
lstm_model = Sequential([
    LSTM(64, return_sequences=True, input_shape=(n_timesteps, 1)),
    Dropout(0.2),
    LSTM(64),
    Dropout(0.2),
    Dense(64, activation="relu"),
    Dense(num_classes, activation="softmax") # num_classes output classes
])

# Compiling model
lstm_model.compile(optimizer="adam", loss="categorical_crossentropy",
    metrics=["accuracy"])

# Set up EarlyStopping

```

```

early_stop = EarlyStopping(monitor='val_loss', patience=5,
    ↪restore_best_weights=True)

# Training the LSTM model
history = lstm_model.fit(
    X_train_lstm,
    y_train_cat,
    validation_data=(X_val_lstm, y_val_cat),
    epochs=20,
    batch_size=32,
    callbacks=[early_stop]
)

```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/rnn/rnn.py:200:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first
layer in the model instead.

```
super().__init__(**kwargs)
```

Epoch 1/20

1885/1885 409s 214ms/step

- accuracy: 0.8662 - loss: 0.5926 - val_accuracy: 0.8702 - val_loss: 0.5233

Epoch 2/20

1885/1885 457s 222ms/step

- accuracy: 0.8731 - loss: 0.5082 - val_accuracy: 0.9118 - val_loss: 0.3491

Epoch 3/20

1885/1885 417s 221ms/step

- accuracy: 0.9116 - loss: 0.3514 - val_accuracy: 0.9157 - val_loss: 0.3272

Epoch 4/20

1885/1885 442s 222ms/step

- accuracy: 0.9158 - loss: 0.3330 - val_accuracy: 0.9166 - val_loss: 0.3230

Epoch 5/20

1885/1885 444s 222ms/step

- accuracy: 0.9168 - loss: 0.3235 - val_accuracy: 0.9201 - val_loss: 0.3204

Epoch 6/20

1885/1885 443s 223ms/step

- accuracy: 0.9238 - loss: 0.3121 - val_accuracy: 0.9255 - val_loss: 0.2999

Epoch 7/20

1885/1885 445s 225ms/step

- accuracy: 0.9255 - loss: 0.3114 - val_accuracy: 0.9257 - val_loss: 0.3103

Epoch 8/20

1885/1885 439s 223ms/step

- accuracy: 0.9276 - loss: 0.3021 - val_accuracy: 0.9276 - val_loss: 0.2937

Epoch 9/20

1885/1885 440s 222ms/step

- accuracy: 0.9290 - loss: 0.2934 - val_accuracy: 0.8702 - val_loss: 0.5521

Epoch 10/20

1885/1885 440s 221ms/step

```

- accuracy: 0.8705 - loss: 0.5454 - val_accuracy: 0.8702 - val_loss: 0.5293
Epoch 11/20
1885/1885          442s 221ms/step
- accuracy: 0.8795 - loss: 0.4843 - val_accuracy: 0.9273 - val_loss: 0.2882
Epoch 12/20
1885/1885          444s 222ms/step
- accuracy: 0.9399 - loss: 0.2484 - val_accuracy: 0.9491 - val_loss: 0.2041
Epoch 13/20
1885/1885          427s 214ms/step
- accuracy: 0.9500 - loss: 0.2098 - val_accuracy: 0.9532 - val_loss: 0.1797
Epoch 14/20
1885/1885          456s 222ms/step
- accuracy: 0.9540 - loss: 0.1767 - val_accuracy: 0.9548 - val_loss: 0.1637
Epoch 15/20
1885/1885          427s 214ms/step
- accuracy: 0.9600 - loss: 0.1543 - val_accuracy: 0.9532 - val_loss: 0.1636
Epoch 16/20
1885/1885          459s 223ms/step
- accuracy: 0.9613 - loss: 0.1477 - val_accuracy: 0.9629 - val_loss: 0.1353
Epoch 17/20
1885/1885          439s 221ms/step
- accuracy: 0.9638 - loss: 0.1391 - val_accuracy: 0.9665 - val_loss: 0.1353
Epoch 18/20
1885/1885          442s 221ms/step
- accuracy: 0.9655 - loss: 0.1346 - val_accuracy: 0.9664 - val_loss: 0.1409
Epoch 19/20
1885/1885          442s 222ms/step
- accuracy: 0.9646 - loss: 0.1381 - val_accuracy: 0.9684 - val_loss: 0.1201
Epoch 20/20
1885/1885          442s 221ms/step
- accuracy: 0.9706 - loss: 0.1173 - val_accuracy: 0.9692 - val_loss: 0.1141

```

LSTM history

```

[ ]: plt.figure(figsize=(12, 5))

# Loss plot
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('LSTM Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

# Accuracy plot
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Train Accuracy')

```

```

plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('LSTM Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.tight_layout()
plt.show()

# Extract final training and validation accuracy and loss values
final_train_accuracy = history.history['accuracy'][-1] * 100
final_val_accuracy = history.history['val_accuracy'][-1] * 100

final_train_loss = history.history['loss'][-1]
final_val_loss = history.history['val_loss'][-1]

final_train_loss_pct = final_train_loss * 100
final_val_loss_pct = final_val_loss * 100


# Create a summary DataFrame for Accuracy
accuracy_summary = pd.DataFrame({
    'Phase': ['Training Accuracy', 'Validation Accuracy'],
    'Accuracy (%)': [final_train_accuracy, final_val_accuracy]
})

print("Final Accuracy Summary:")
display(accuracy_summary)

# Create a summary DataFrame for Loss
loss_summary = pd.DataFrame({
    'Phase': ['Training Loss', 'Validation Loss'],
    'Loss Value': [final_train_loss, final_val_loss],
    'Loss (%)': [final_train_loss_pct, final_val_loss_pct]
})

print("Final Loss Summary:")
display(loss_summary)

```



Final Accuracy Summary:

	Phase	Accuracy (%)
0	Training Accuracy	96.997380
1	Validation Accuracy	96.922874

Final Loss Summary:

	Phase	Loss Value	Loss (%)
0	Training Loss	0.118750	11.875039
1	Validation Loss	0.114145	11.414526

Defining Top-k Accuracy Function

```
[ ]: def top_k_accuracy(y_true, y_prob, k=5):
    top_k_preds = np.argsort(y_prob, axis=1)[: , -k:]
    return np.mean([1 if y_true.values[i] in top_k_preds[i] else 0 for i in
    ↪range(len(y_true))])
```

LR

```
[ ]: # --- Logistic Regression Metrics Evaluation and Display as DataFrames ---

# Predictions and predicted probabilities
y_val_pred_lr = grid_lr.predict(X_val_split)
y_val_prob_lr = grid_lr.predict_proba(X_val_split)

# Compute the classification report
lr_report = classification_report(y_val_split, y_val_pred_lr, output_dict=True)
lr_report_df = pd.DataFrame(lr_report).transpose()
print("Logistic Regression Classification Report:")
display(lr_report_df)

# Compute overall accuracy
```

```

accuracy_lr = accuracy_score(y_val_split, y_val_pred_lr)

classes = np.unique(y_val_split)
y_val_onehot = label_binarize(y_val_split, classes=classes)

# Compute ROC AUC Score (macro average, one-vs-rest)
roc_auc_lr = roc_auc_score(y_val_onehot, y_val_prob_lr, average="macro",
    ↪multi_class="ovr")

# Compute AUPR for each class and then macro average
aupr_scores_lr = []
for i in range(len(classes)):
    precision, recall, _ = precision_recall_curve(y_val_onehot[:, i],
    ↪y_val_prob_lr[:, i])

    sorted_indices = np.argsort(recall)
    recall_sorted = recall[sorted_indices]
    precision_sorted = precision[sorted_indices]
    aupr = auc(recall_sorted, precision_sorted)
    aupr_scores_lr.append(aupr)
macro_aupr_lr = np.mean(aupr_scores_lr)

# Compute Top-5 Accuracy (mAP)
top5_map_lr = top_k_accuracy(y_val_split, y_val_prob_lr, k=5)

summary_lr = pd.DataFrame({
    'Metric': ['Accuracy', 'ROC AUC (macro, OVR)', 'Macro AUPR', 'Top-5 mAP'],
    'Value': [accuracy_lr, roc_auc_lr, macro_aupr_lr, top5_map_lr]
})
print("Overall Logistic Regression Performance Metrics:")
display(summary_lr)

aupr_scores_lr_df = pd.DataFrame({
    'Class': classes,
    'AUPR': aupr_scores_lr
})
print("AUPR Scores for Each Class:")
display(aupr_scores_lr_df)

```

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0.0	0.958226	0.992913	0.975261	13122.000000
1.0	0.882845	0.541026	0.670906	390.000000
2.0	0.858491	0.610738	0.713725	596.000000
3.0	0.459016	0.250000	0.323699	112.000000
4.0	0.974934	0.860303	0.914038	859.000000

accuracy	0.953047	0.953047	0.953047	0.953047
macro avg	0.826702	0.650996	0.719526	15079.000000
weighted avg	0.949578	0.953047	0.948725	15079.000000

Overall Logistic Regression Performance Metrics:

	Metric	Value
0	Accuracy	0.953047
1	ROC AUC (macro, OVR)	0.945421
2	Macro AUPR	0.743697
3	Top-5 mAP	1.000000

AUPR Scores for Each Class:

	Class	AUPR
0	0.0	0.983444
1	1.0	0.621228
2	2.0	0.786060
3	3.0	0.392902
4	4.0	0.934850

RF

```
[ ]: # --- Random Forest Metrics Evaluation and Display as DataFrames ---

# Predictions and predicted probabilities for Random Forest
y_val_pred_rf = rf_model.predict(X_val_split)
y_val_prob_rf = rf_model.predict_proba(X_val_split)

# Generate classification report as a dictionary
rf_report = classification_report(y_val_split, y_val_pred_rf, output_dict=True)
rf_report_df = pd.DataFrame(rf_report).transpose()
print("\nRandom Forest Classification Report:")
display(rf_report_df)

# Compute overall accuracy
accuracy_rf = accuracy_score(y_val_split, y_val_pred_rf)

# ROC AUC Score
roc_auc_rf = roc_auc_score(y_val_onehot, y_val_prob_rf, average="macro",
    multi_class="ovr")

# Compute AUPR for each class with sorted recall for proper AUC calculation
aupr_scores_rf = []
for i in range(len(classes)):
    precision, recall, _ = precision_recall_curve(y_val_onehot[:, i],
        y_val_prob_rf[:, i])
    sorted_indices = np.argsort(recall)
    recall_sorted = recall[sorted_indices]
    precision_sorted = precision[sorted_indices]
```

```

    auapr = auc(recall_sorted, precision_sorted)
    auapr_scores_rf.append(auapr)
macro_auapr_rf = np.mean(auapr_scores_rf)

# Compute Top-5 mAP (top-5 accuracy)
top5_map_rf = top_k_accuracy(y_val_split, y_val_prob_rf, k=5)

# Create a summary DataFrame for overall metrics
summary_rf = pd.DataFrame({
    'Metric': ['Accuracy', 'ROC AUC (macro, OVR)', 'Macro AUPR', 'Top-5 mAP'],
    'Value': [accuracy_rf, roc_auc_rf, macro_auapr_rf, top5_map_rf]
})
print("Overall Random Forest Performance Metrics:")
display(summary_rf)

# Create a DataFrame for AUPR scores per class
auapr_scores_rf_df = pd.DataFrame({
    'Class': classes,
    'AUPR': auapr_scores_rf
})
print("Random Forest AUPR Scores for Each Class:")
display(auapr_scores_rf_df)

```

Random Forest Classification Report:

	precision	recall	f1-score	support
0.0	0.976162	0.998628	0.987267	13122.000000
1.0	0.976285	0.633333	0.768274	390.000000
2.0	0.978389	0.835570	0.901357	596.000000
3.0	0.811765	0.616071	0.700508	112.000000
4.0	0.996287	0.937136	0.965807	859.000000
accuracy	0.976391	0.976391	0.976391	0.976391
macro avg	0.947778	0.804148	0.864643	15079.000000
weighted avg	0.976179	0.976391	0.974855	15079.000000

Overall Random Forest Performance Metrics:

	Metric	Value
0	Accuracy	0.976391
1	ROC AUC (macro, OVR)	0.993818
2	Macro AUPR	0.932766
3	Top-5 mAP	1.000000

Random Forest AUPR Scores for Each Class:

	Class	AUPR
0	0.0	0.999095
1	1.0	0.897132
2	2.0	0.976879

```
3    3.0  0.795239
4    4.0  0.995485
```

LSTM

```
[ ]: # --- LSTM Model Metrics Evaluation and Display as DataFrames ---

# Get predictions and predicted probabilities for LSTM
y_val_pred_lstm_prob = lstm_model.predict(X_val_lstm)
y_val_pred_lstm = np.argmax(y_val_pred_lstm_prob, axis=1)

# Generate classification report and convert it to a DataFrame
print("\nLSTM Model Classification Report:")
lstm_report = classification_report(y_val_split, y_val_pred_lstm,
    ↳output_dict=True)
lstm_report_df = pd.DataFrame(lstm_report).transpose()
display(lstm_report_df)

# Calculate overall accuracy
accuracy_lstm = accuracy_score(y_val_split, y_val_pred_lstm)
print("LSTM Model Accuracy:", accuracy_lstm)

# Compute ROC AUC Score
roc_auc_lstm = roc_auc_score(y_val_onehot, y_val_pred_lstm_prob,
    ↳average="macro", multi_class="ovr")
print("LSTM Model ROC AUC (macro, OVR):", roc_auc_lstm)

# Compute AUPR for each class
aupr_scores_lstm = []
for i in range(len(classes)):
    precision, recall, _ = precision_recall_curve(y_val_onehot[:, i],
    ↳y_val_pred_lstm_prob[:, i])
    sorted_indices = np.argsort(recall)
    recall_sorted = recall[sorted_indices]
    precision_sorted = precision[sorted_indices]
    aupr = auc(recall_sorted, precision_sorted)
    aupr_scores_lstm.append(aupr)

macro_aupr_lstm = np.mean(aupr_scores_lstm)
print("LSTM Model Macro AUPR:", macro_aupr_lstm)

# Compute Top-5 Accuracy (mAP)
top5_map_lstm = top_k_accuracy(y_val_split, y_val_pred_lstm_prob, k=5)
print("LSTM Model Top-5 mAP:", top5_map_lstm)

#summary DataFrame for overall LSTM metrics
summary_lstm = pd.DataFrame({
    'Metric': ['Accuracy', 'ROC AUC (macro, OVR)', 'Macro AUPR', 'Top-5 mAP'],
```

```

        'Value': [accuracy_lstm, roc_auc_lstm, macro_aupr_lstm, top5_map_lstm]
    })
    print("Overall LSTM Performance Metrics:")
    display(summary_lstm)

    #AUPR scores for each class
    aupr_scores_lstm_df = pd.DataFrame({
        'Class': classes,
        'AUPR': aupr_scores_lstm
    })
    print("LSTM AUPR Scores for Each Class:")
    display(aupr_scores_lstm_df)

```

472/472 34s 73ms/step

LSTM Model Classification Report:

	precision	recall	f1-score	support
0.0	0.979641	0.990093	0.984839	13122.000000
1.0	0.771341	0.648718	0.704735	390.000000
2.0	0.899461	0.840604	0.869037	596.000000
3.0	0.617188	0.705357	0.658333	112.000000
4.0	0.982587	0.919674	0.950090	859.000000
accuracy	0.969229	0.969229	0.969229	0.969229
macro avg	0.850044	0.820889	0.833407	15079.000000
weighted avg	0.968560	0.969229	0.968613	15079.000000

LSTM Model Accuracy: 0.9692287286955369

LSTM Model ROC AUC (macro, OVR): 0.9766423819858623

LSTM Model Macro AUPR: 0.8666640873982402

LSTM Model Top-5 mAP: 1.0

Overall LSTM Performance Metrics:

	Metric	Value
0	Accuracy	0.969229
1	ROC AUC (macro, OVR)	0.976642
2	Macro AUPR	0.866664
3	Top-5 mAP	1.000000

LSTM AUPR Scores for Each Class:

	Class	AUPR
0	0.0	0.995887
1	1.0	0.728098
2	2.0	0.944513
3	3.0	0.691998
4	4.0	0.972825

ROC curves

```

[ ]: sns.set_style("whitegrid")

rows = (n_classes + 1) // 2
cols = 2

fig, axes = plt.subplots(rows, cols, figsize=(12, 6 * rows))
axes = axes.flatten()

for i in range(n_classes):
    fpr_lr, tpr_lr, _ = roc_curve(y_val_onehot[:, i], y_val_prob_lr[:, i])
    fpr_rf, tpr_rf, _ = roc_curve(y_val_onehot[:, i], y_val_prob_rf[:, i])
    fpr_lstm, tpr_lstm, _ = roc_curve(y_val_onehot[:, i], y_val_pred_lstm_prob[:, i])

    ax = axes[i]
    ax.plot(fpr_lr, tpr_lr, label=f'LR AUC: {auc(fpr_lr, tpr_lr):.2f}',
            color='blue', linewidth=2.5)
    ax.plot(fpr_rf, tpr_rf, label=f'RF AUC: {auc(fpr_rf, tpr_rf):.2f}',
            color='green', linewidth=2.5)
    ax.plot(fpr_lstm, tpr_lstm, label=f'LSTM AUC: {auc(fpr_lstm, tpr_lstm):.2f}',
            color='red', linewidth=2.5)
    ax.plot([0, 1], [0, 1], 'k--', linewidth=1) # Diagonal reference line

    ax.set_xlim([0, 1])
    ax.set_ylim([0, 1])
    ax.set_xticks([0, 0.2, 0.4, 0.6, 0.8, 1])
    ax.set_yticks([0, 0.2, 0.4, 0.6, 0.8, 1])
    ax.set_title(f'ROC Curve - Class {i}', fontsize=14, fontweight='bold')
    ax.set_xlabel('False Positive Rate', fontsize=12)
    ax.set_ylabel('True Positive Rate', fontsize=12)
    ax.legend(loc='lower right', fontsize=11)

for j in range(n_classes, len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

```

